

EU AI Act Compliance Audit Report

Automated pipeline output | Timestamp: 23/05/2026, 7.01.24

Assessment Metadata

- Session ID: session_nb323forz
- Compliance Score: 10/100
- Regulatory Status: ACTION REQUIRED

1. Extracted Use-Case Facts (Source: Input Parser Agent)

Core Purpose: Unknown purpose.

Target Sector: General AI

2. EU AI Act Path Classifications & Citations (Source: Decision Tree)

Unclassified / Out of Scope (Article 2)

EU AI Act Flowchart Citations:

- Article 2 (Scope & Exclusions check - Determines if the system's deployment affects the EU market)
- Chapter IX (Post-Market Surveillance - Conformance framework for monitoring, info sharing, and market oversight)
- Article 99 (General Penalties and Fines framework - Outlines overall administrative enforcement)
- EU Legal Supremacy Disclaimer (The AI Act operates alongside, and does not replace, other sectoral EU laws)

3. Deduced Governance Observations (Source: Judge of Governance)

Lifecycle Logging & Documentation: Default document assessment notes.

4. Risk Boundaries: Assumptions & Gaps Checkers

Operational & Technical Assumptions Checked:

- **Operational & Technical:** The proposal for a "Unknown purpose." targeting end users and system operators assumes high-fidelity "Human-in-the-loop" oversight, claiming: "system-level supervision". However, in real-world workflows, this creates a dangerous assumption that operators will actively review and verify the generated "automated decision outputs" before taking action. In high-throughput, resource-scarce environments, clinical or corporate operators are highly prone to "automation bias"—passively accepting or rubber-stamping recommendations without independent cognitive verification, effectively transforming an advisory tool into an autonomous pipeline.
- **Supply Chain & Role:** The proposal for a "Unknown purpose." within the General AI sector does not explicitly specify reliance on upstream GPAI models. This assumes the system is built entirely on custom-trained, proprietary models, creating a clean boundary as the primary "Provider". However, this assumes the developers possess the exhaustive in-house infrastructure and data validation pipelines required to verify model weights, bias controls, and alignment criteria without relying on third-party Service Level Agreements (SLAs).
- **Legal Interpretation:** The proposal for a "Unknown purpose." within the General AI assumes that its risk tier is determined solely by the marketing intent of the proposal. However, under the EU AI Act, system classification is based on actual operational use cases and risk domains, which must be carefully audited to ensure they do not cross into high-risk categories through secondary feature bloat. Given the potential impact of "general human actions and livelihoods", a failure to preemptively align with statutory High-Risk obligations risks immediate market withdrawal or massive administrative fines.
- **Vulnerability & Impact (Optional):** NOT APPLICABLE. The system does not target or directly impact human safety, livelihoods, or vulnerable populations. No dangerous vulnerability exploitation or asymmetrical power assumptions are present.
- **Temporal & Lifecycle:** Deploying the "Unknown purpose." in the context of "production and market deployment" assumes that compliance obligations only commence once the product is fully commercialized. However, the EU AI Act applies the moment a system is 'placed on the market' or 'put into service' (even for trial/pilot programs in "production and market deployment"). The proposal assumes its rollout is exempt under research/sandbox clauses, but any operational use with live data activates full legal liabilities immediately. Furthermore, any future model updates or weight adjustments in "production and market deployment" are assumed to be simple maintenance patches, whereas they actually constitute a 'substantial modification' requiring brand-new conformity certifications.
- **Geographic & Jurisdictional:** The deployment plan for "Unknown purpose." via "production and market deployment" assumes that geographic boundaries or external hosting provide a shield against EU regulatory jurisdiction. However, Article 2 of the AI Act establishes strict extraterritoriality: if the system's outputs are used or have effects within the European Union, the system is fully subject to the AI Act. This assumes that the location of the servers, developers, or operators in "production and market

deployment" exempts the project, which is a major legal liability if EU residents are affected.

- Data Provenance: The proposal for a "Unknown purpose." assumes that its input datasets ("system input datasets") are inherently balanced, high-quality, and legally compliant for processing. However, training or running models on "system input datasets" assumes the absence of systemic or historical biases. In practice, this data frequently reflects entrenched patterns that, when processed, mathematically codify and perpetuate bias. This violates Article 10's strict data governance requirements, especially considering the system's potential impact on "general human actions and livelihoods".

Identified Regulatory Assessment Gaps:

- [GAP IN ASSESSMENT] Missing Governance Area: 'Risk Management'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Transparency'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Human Oversight'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Monitoring'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Logging'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Accountability'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.
- [GAP IN ASSESSMENT] Missing Governance Area: 'Role Clarity'. No operational notes, procedural evidence, or safety strategies were provided in the proposal for this AI Act mandate. This creates high liability and audit compliance risks.

5. Confidence Prevention & Auditing Queries (Prevention of Confidence Agent)

AI Compliance Expert Auditing Queries:

1. ****How will the system dynamically identify, version, and log model performance and telemetry deviations when underlying neural network weights or system prompts are updated?****
2. ****What rigorous, quantitative verification benchmarks (e.g., safety, accuracy, robust adversarial testing) will be used to certify the security posture of this AI deployment before launching in live production contexts?****

Advisory Footnote: This report displays a multi-agent automated compliance assessment to facilitate regulatory review. It does not constitute formal legal counsel or binding certification under the EU AI Act.